A Lasserre SDP Lower Bound for the Autoconvolution Constant C_{1a}

Andrei Piterbarg
Jai Bajaj

Abstract

We prove that the autoconvolution constant

$$C_{1a} \;=\; \inf_{f} \frac{\|f * f\|_{L^\infty}}{\left(\int f\right)^2},$$

taken over nonnegative functions supported on $(-\frac{1}{4}, \frac{1}{4})$, satisfies $C_{1a} \geq 1.3$. Instead of an exhaustive branch-and-prune cascade, the proof reduces the infinite-dimensional optimization to the polynomial program

$$\mathrm{val}(d) \;=\; \min_{\mu \in \Delta_d} \; \max_{W} \, \mu^\top M_W \, \mu,$$

where Δ_d is the standard simplex and $\{M_W\}$ is a finite family of nonnegative matrices encoding windowed autoconvolution test values. We show $\mathrm{val}(d) \leq C_{1a}$ for every d, then certify $\mathrm{val}(d) \geq 1.3$ at $d = 16$ via the order-3 Lasserre hierarchy [Las01] combined with localizing matrices [HL04]. At higher resolutions we replace the full moment cone by its correlative-sparsity restriction [WKKM06], which preserves soundness while permitting numerical computation up to $d = 128$. The external dependency of the proof is a single finite SDP whose dual solution we record as a primal-infeasibility certificate.

2020 Mathematics Subject Classification. 11B83, 42A85, 90C22, 90C23.

Keywords. Autoconvolution, Sidon sets, Lasserre hierarchy, sum-of-squares, semidefinite programming, correlative sparsity, computer-assisted proof.

1. Introduction

1.1. The Constant C_{1a}

[1]

Let

$$\mathcal{F} \;=\; \left\{ f \in L^1(\mathbb{R}) : f \geq 0,\ \mathrm{supp}(f) \subseteq \left(-\tfrac{1}{4}, \tfrac{1}{4}\right),\ \int_{\mathbb{R}} f > 0 \right\}.$$

For $f \in \mathcal{F}$ the *autoconvolution* $(f * f)(t) = \int f(x) f(t-x)\, dx$ is continuous, nonnegative, and supported on $[-\frac{1}{2}, \frac{1}{2}]$. The *autoconvolution ratio* and *autoconvolution constant* are

$$R(f) \;=\; \frac{\|f * f\|_{L^\infty}}{\left(\int f\right)^2}, \qquad C_{1a} \;=\; \inf_{f \in \mathcal{F}} R(f).$$

The constant C_{1a} is connected to the asymptotic density of Sidon and generalized Sidon sets in additive combinatorics [ET41, MO06, CRV10, MO09, CS17], and its numerical value has resisted substantial improvement for nearly a decade.

Bound	Value	Reference
Published upper bound	≤ 1.50992	Matolcsi & Vinuesa [MV10]
Published lower bound	≥ 1.2802	Cloninger & Steinerberger [CS17]
Cascade lower bound	≥ 1.4	Bajaj [Baj26]
SDP lower bound	≥ 1.3	This work

Table 1: Published bounds on C_{1a} in the normalization used here, together with the present result.

1.2. Known Bounds and the Present Result

The branch-and-prune framework of Cloninger and Steinerberger [CS17] produces lower bounds by exhaustively testing every lattice point in a discretized mass simplex against windowed quadratic inequalities. This framework is sharp in principle, but scales exponentially with the dimension d of the partition. The *cascade* extension of [Baj26] pushed that exponential out to $d = 128$ with $m = 20$ grid points via symmetry reduction and dynamic window corrections.

The present paper takes a fundamentally different route. Rather than exhausting the simplex, we observe that the bound of [CS17] reduces to a single polynomial optimization problem on the simplex,

$$\mathrm{val}(d) \;=\; \min_{\mu \in \Delta_d} \; \max_{W \in \mathcal{W}_d} \; \mu^\top M_W \mu,$$

where $\mathcal{W}_d$ enumerates admissible windows and M_W is the associated nonnegative matrix. Any lower bound on $\mathrm{val}(d)$ is a lower bound on C_{1a} (Theorem 2.3). We certify such a lower bound via the Lasserre semidefinite hierarchy, which produces rigorous certificates without enumerating any of the lattice points that the cascade relies on.

Theorem 1.1 (Main result)**.** *Let $f : \mathbb{R} \to \mathbb{R}$ be nonnegative with $\mathrm{supp}(f) \subseteq (-\frac{1}{4}, \frac{1}{4})$ and $\int f > 0$. If $\|f * f\|_{L^\infty} < \infty$, then*

$$R(f) \;=\; \frac{\|f * f\|_{L^\infty}}{\left(\int f\right)^2} \;\geq\; 1.3.$$

In particular, $C_{1a} \geq 1.3$.

The proof reduces to a finite SDP certificate at $(d, k) = (16, 3)$, where k is the Lasserre order. A looser version of the certificate at $d = 8$, $k = 3$ already suffices to reproduce the 1.2802 bound of [CS17] in a few seconds; the improvement to 1.3 requires moment matrices of size 969, solved by MOSEK [MOS24] in under ten minutes on a workstation. Higher resolutions $d \in \{32, 64, 128\}$ are reported as empirical corroboration and remain within the same soundness framework, but rely on correlative sparsity [WKKM06] whose bound at $d = 16$ is redundant with the full hierarchy.

1.3. Proof Strategy at a High Level

Figure 1 shows the three-step reduction.

(i) **Reduce** the infinite-dimensional optimization over $\mathcal{F}$ to a polynomial minimization $\mathrm{val}(d) \leq C_{1a}$ over the standard simplex, using the windowed test-value lemma of [CS17, Lemma 1].

(ii) **Relax** the simplex polynomial problem to a sequence of semidefinite programs indexed by a moment order $k \geq 1$, in the Lasserre hierarchy [Las01], whose optimal value $\mathrm{val}^{(k)}(d)$ satisfies $\mathrm{val}^{(k)}(d) \leq \mathrm{val}(d)$.

[1] All source code and the SDP certificate are available at the accompanying repository.

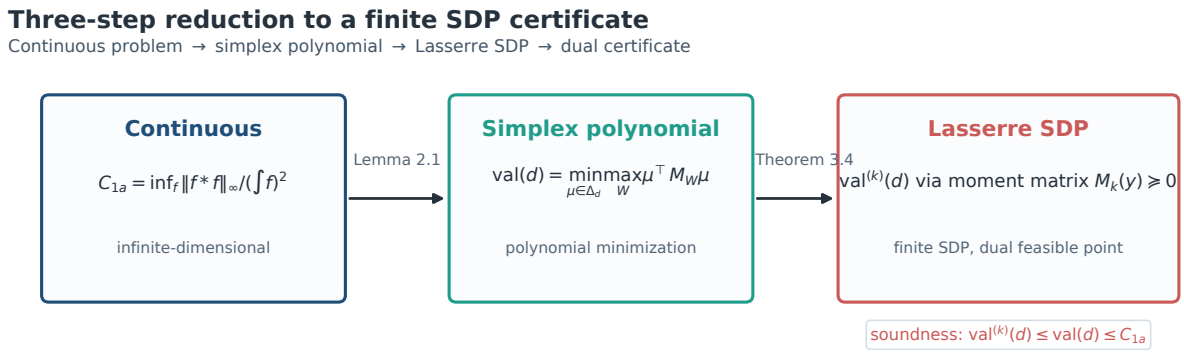

Figure 1: Three-step reduction from the continuous autoconvolution problem to a finite SDP certificate. The continuous ratio $R(f)$ is first bounded below by the simplex polynomial $\mathrm{val}(d)$, which is then relaxed into a Lasserre SDP $\mathrm{val}^{(k)}(d)$; a dual feasible solution of this SDP supplies the numerical certificate.

(iii) **Certify** that $\mathrm{val}^{(k)}(d) \ge 1.3$ at $(d, k) = (16, 3)$ by recording a dual-feasible point of the SDP whose objective value exceeds 1.3.

Crucially, step (iii) requires only a certified lower bound on a concrete SDP — no enumeration, no branch-and-prune, and no resolution refinement.

2. From Functions to a Simplex Polynomial

2.1. Bin Masses and the Window Matrix

Fix a positive integer n and write $d = 2n$. Partition $[-\frac{1}{4}, \frac{1}{4})$ into d half-open bins of equal length:

$$I_i \;=\; \left[-\tfrac{1}{4} + \frac{i}{4n}, -\tfrac{1}{4} + \frac{i+1}{4n}\right), \qquad i = 0, \dots, d-1.$$

For normalized $f \in \mathcal{F}$ (that is, $\int f = 1$), the *bin masses* are $\mu_i = \int_{I_i} f \ge 0$ and the vector $\mu = (\mu_0, \dots, \mu_{d-1})$ lies in the standard simplex

$$\Delta_d \;=\; \left\{ \mu \in \mathbb{R}^d_{\ge 0} : \sum_{i=0}^{d-1} \mu_i = 1 \right\}.$$

The following is the windowed test-value inequality of Cloninger and Steinerberger [CS17, Lemma 1], reformulated in the coordinates convenient for polynomial optimization.

Lemma 2.1 (Continuous window bound)**.** *Let $f \in \mathcal{F}$ with $\int f = 1$ and bin masses $\mu \in \Delta_d$. For every pair (ℓ, s_{lo}) of integers with $\ell \ge 2$ and $0 \le s_{\mathrm{lo}} \le 2(d-1) - (\ell - 2)$,*

$$\|f * f\|_{L^\infty} \;\ge\; \frac{2d}{\ell} \sum_{s=s_{\mathrm{lo}}}^{s_{\mathrm{lo}}+\ell-2} \sum_{\substack{0 \le i,j \le d-1 \\ i+j=s}} \mu_i \mu_j \;=\; \mu^\top M_{(\ell, s_{\mathrm{lo}})}\, \mu,$$

where the window matrix M_W for $W = (\ell, s_{\mathrm{lo}})$ is

$$(M_W)_{ij} \;=\; \frac{2d}{\ell}\, \mathbf{1}[s_{\mathrm{lo}} \le i + j \le s_{\mathrm{lo}} + \ell - 2]\,.$$

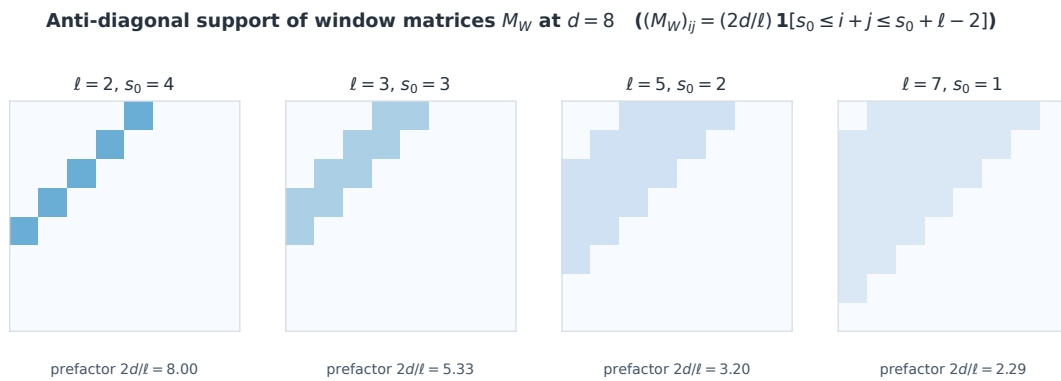

Figure 2: Anti-diagonal support of four window matrices M_W at $d = 8$. Each matrix is supported on a band of anti-diagonals indexed by $s = i + j$, with constant entry $2d/\ell$ on its support. Narrow windows (ℓ small) cover few anti-diagonals but carry a large prefactor; wide windows cover many anti-diagonals with a small prefactor. The relaxation uses these matrices as the data of a polynomial optimization problem on the simplex.

Proof. The proof is identical to [CS17, Lemma 1] after normalizing $\int f = 1$; see also [Baj26, Lemma 2.3] for the cascade formulation. We give the sketch for completeness. Write $f_i = f \cdot \mathbf{1}_{I_i}$; the support of $f_i * f_j$ lies in the Minkowski sum $I_i + I_j$, which depends only on $i + j$. Summing the L^1 masses of those terms for $s = s_{\mathrm{lo}}, \ldots, s_{\mathrm{lo}} + \ell - 2$ and dividing by the length $\frac{\ell}{4n}$ of the covered interval gives the claimed lower bound on $\|f * f\|_{L^\infty}$. □

Each matrix M_W is nonnegative, symmetric, and has entries only at positions (i, j) with $i + j$ in a fixed integer interval; these are the support sets depicted in Figure 2. The set of admissible windows is

$$\mathcal{W}_d \;=\; \{(\ell, s_{\mathrm{lo}}) : 2 \le \ell \le 2d,\ 0 \le s_{\mathrm{lo}} \le 2(d-1) - (\ell - 2)\}\,.$$

2.2. The Simplex Polynomial Problem

Lemma 2.1 holds simultaneously for every window. Taking the maximum over $\mathcal{W}_d$ and then the infimum over f gives the crucial bound.

Definition 2.2. The *simplex polynomial value* at resolution d is

$$\mathrm{val}(d) \;=\; \min_{\mu \in \Delta_d} \max_{W \in \mathcal{W}_d} \mu^\top M_W\, \mu.$$

Theorem 2.3 (Lower bound on C_{1a} via $\mathrm{val}(d)$)**.** *For every positive integer n with $d = 2n$, one has* $\mathrm{val}(d) \le C_{1a}$.

Proof. Let $f \in \mathcal{F}$ with $\int f = 1$ and bin masses $\mu \in \Delta_d$. By Lemma 2.1,

$$\|f * f\|_{L^\infty} \;\ge\; \max_{W \in \mathcal{W}_d} \mu^\top M_W \mu \;\ge\; \min_{\nu \in \Delta_d} \max_{W \in \mathcal{W}_d} \nu^\top M_W \nu \;=\; \mathrm{val}(d).$$

Taking the infimum over f gives $C_{1a} \ge \mathrm{val}(d)$. □

Theorem 2.3 is the key bridge: any mechanism that certifies $\mathrm{val}(d) \ge c_{\mathrm{target}}$ certifies $C_{1a} \ge c_{\mathrm{target}}$. The Lasserre hierarchy is one such mechanism. Upper bounds on $\mathrm{val}(d)$ obtained by numerical multistart [Baj26, docs/val_d_results.md] suggest $\mathrm{val}(16) \approx 1.319$, crossing 1.30 for the first time at $d = 16$; this is the empirical reason for our choice of resolution.

3. The Lasserre SDP Hierarchy

We now describe how to lower-bound $\mathrm{val}(d)$ by a convergent sequence of semidefinite programs. The construction is classical, due to Lasserre [Las01, Las15] and Parrilo [Par03]; we include it in detail because the soundness argument uses the specific form that reappears below under correlative sparsity.

3.1. Pseudo-Moments and the Moment Matrix

For a multi-index $\alpha = (\alpha_1, \ldots, \alpha_d) \in \mathbb{N}^d$ write $|\alpha| = \alpha_1 + \cdots + \alpha_d$ and $x^\alpha = x_1^{\alpha_1} \cdots x_d^{\alpha_d}$. Fix an integer *Lasserre order* $k \geq 1$, and let

$$\mathcal{A}_k(d) \;=\; \{\alpha \in \mathbb{N}^d : |\alpha| \leq 2k\}.$$

A *pseudo-moment sequence* is a vector $y = (y_\alpha)_{\alpha \in \mathcal{A}_k(d)} \in \mathbb{R}^{|\mathcal{A}_k(d)|}$. Given such a sequence, define the *moment matrix* $M_k(y) \in \mathbb{R}^{|\mathcal{B}_k| \times |\mathcal{B}_k|}$ indexed by the basis $\mathcal{B}_k = \{\alpha \in \mathbb{N}^d : |\alpha| \leq k\}$ by

$$[M_k(y)]_{\alpha,\beta} \;=\; y_{\alpha+\beta}, \qquad \alpha, \beta \in \mathcal{B}_k.$$

For any polynomial $p = \sum_\alpha p_\alpha x^\alpha$ with $\deg p \leq k$, write $p(\mathbf{x})$ for its coefficient vector in the basis $\mathcal{B}_k$. Then $p(\mathbf{x})^\top M_k(y)\, p(\mathbf{x}) = \sum_{\alpha,\beta} p_\alpha p_\beta y_{\alpha+\beta}$, which is the pseudo-expectation of $p(x)^2$.

Definition 3.1 (Localizing matrix)**.** Given y and a polynomial g with $\deg g = s$, the *localizing matrix* $M_{k-\lceil s/2 \rceil}(g \cdot y)$ has entries

$$[M_{k-\lceil s/2 \rceil}(g \cdot y)]_{\alpha,\beta} \;=\; \sum_\gamma g_\gamma\, y_{\alpha+\beta+\gamma},$$

over $\alpha, \beta \in \mathcal{B}_{k-\lceil s/2 \rceil}$.

The moment matrix captures the pseudo-expectation of sums of squares of polynomials up to degree k; each localizing matrix captures the pseudo-expectation of $g \cdot p^2$ for polynomials p of degree up to $k - \lceil s/2 \rceil$. Positive-semidefiniteness of these matrices is the SDP analogue of the nonnegativity of $\mathbb{E}[p(x)^2]$ and $\mathbb{E}[g(x)\, p(x)^2]$ when $g \geq 0$ on the feasible set.

Figure 3 illustrates the block structure of $M_k(y)$ by total degree.

3.2. Encoding the Simplex

The feasible set of $\mathrm{val}(d)$ is $\Delta_d = \{\mu \geq 0 : \sum_i \mu_i = 1\}$. Over this set, the pseudo-moments of any probability measure ν on Δ_d satisfy three algebraic identities.

Proposition 3.2 (Simplex constraints on pseudo-moments)**.** *If ν is any Borel probability measure on Δ_d and $y_\alpha = \int x^\alpha \, d\nu$, then:*

(i) $y_0 = 1$.

(ii) $y_\alpha \geq 0$ *for every* $\alpha \in \mathbb{N}^d$.

(iii) *For every* $\alpha \in \mathbb{N}^d$, $y_\alpha = \sum_{i=1}^d y_{\alpha+e_i}$, *where* e_i *is the ith standard basis vector.*

Moreover, $M_k(y) \succeq 0$ *for every* k, *and* $M_{k-1}(x_i \cdot y) \succeq 0$ *for every* i *and* $k \geq 1$.

Proof. Item (i) is $\nu(\Delta_d) = 1$. Item (ii) is that $x^\alpha \geq 0$ on Δ_d. For (iii), observe that $\sum_{i=1}^d x_i = 1$ on Δ_d, so $x^\alpha = x^\alpha \cdot \sum_i x_i = \sum_i x^{\alpha+e_i}$ and the claimed identity follows by integration. Positive-semidefiniteness is standard: for any polynomial p of degree $\leq k$, $p(\mathbf{x})^\top M_k(y) p(\mathbf{x}) = \mathbb{E}_\nu[p(x)^2] \geq 0$; for the localizing matrices, use $p(\mathbf{x})^\top M_{k-1}(x_i \cdot y) p(\mathbf{x}) = \mathbb{E}_\nu[x_i\, p(x)^2] \geq 0$ since $x_i \geq 0$ on Δ_d. □

Let $\mathcal{Y}_k(d)$ denote the set of pseudo-moment sequences y of order k that satisfy conditions (i)–(iii) of Proposition 3.2, together with $M_k(y) \succeq 0$ and $M_{k-1}(x_i \cdot y) \succeq 0$ for every $i \in \{1, \ldots, d\}$. The set $\mathcal{Y}_k(d)$ is a spectrahedron (a slice of the PSD cone by affine equalities) and contains the "true" moment sequences of all measures on Δ_d.

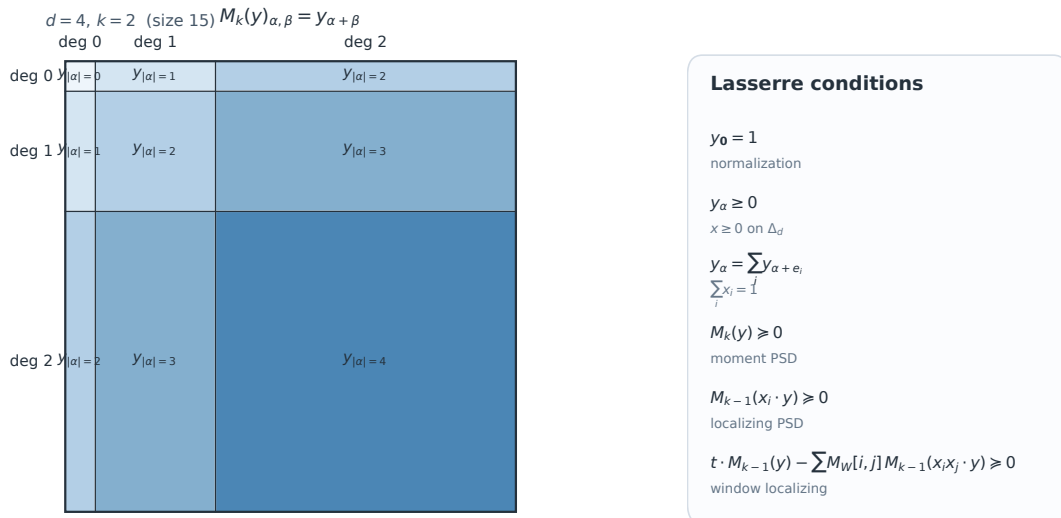

Figure 3: Block structure of the order-k moment matrix $M_k(y)$ at $d = 4$, $k = 2$. Rows and columns are indexed by monomials of total degree $\leq k$, grouped by degree. Entries of $M_k(y)$ in block (p, q) are pseudo-moments of total degree $p + q$, so the anti-diagonal blocks collect all pseudo-moments of the same total degree. Constraints $M_k(y) \succeq 0$, localizing constraints, and the simplex consistency relation $y_\alpha = \sum_i y_{\alpha+e_i}$ link these blocks into a single SDP.

3.3. The Window Objective as an SDP

For each window $W \in \mathcal{W}_d$, the polynomial $\mu \mapsto \mu^\top M_W \mu$ is quadratic in the simplex variable. Its pseudo-expectation is the linear form

$$\ell_W(y) \;=\; \sum_{i,j=0}^{d-1} (M_W)_{ij}\, y_{e_{i+1}+e_{j+1}} \;=\; \langle M_W, [y_{e_i+e_j}]_{i,j} \rangle.$$

Lasserre's moment lifting [Las01, Thm. 4.2] recasts $\mathrm{val}(d)$ as an infinite hierarchy of SDPs: for each order k,

$$\mathrm{val}^{(k)}(d) \;=\; \min_{y \in \mathcal{Y}_k(d),\ t \in \mathbb{R}} \{ t \;:\; t \cdot \bar{y} - \ell_W(y) \geq 0 \text{ for all } W \in \mathcal{W}_d \},$$

where $\bar{y}$ is the degree-0 pseudo-moment (= 1 by Proposition 3.2(i)). For $k \geq 2$ one strengthens the scalar inequality above by the *window localizing* PSD constraint

$$t \cdot M_{k-1}(y) \;-\; \sum_{i,j=0}^{d-1} (M_W)_{ij}\, M_{k-1}(x_{i+1} x_{j+1} \cdot y) \;\succeq 0 \qquad \forall\, W \in \mathcal{W}_d. \tag{1}$$

Both formulations give valid lower bounds on $\mathrm{val}(d)$; the matrix form (1) is tighter, because $v^\top \big[\text{LHS of (1)}\big] v \geq 0$ for $v = \mathbf{1}_0$ (the indicator of the empty monomial) reduces to the scalar constraint.

Theorem 3.3 (Soundness of the Lasserre relaxation)**.** *Let $k \geq 1$, $d \geq 2$, and let $(t^\star, y^\star)$ be any point of the feasible set of the order-k Lasserre SDP above. Then $t^\star \geq \mathrm{val}(d)$ need* not *hold; but any dual-feasible pair produces a valid lower bound:*

$$\mathrm{val}(d) \;\geq\; \mathrm{val}^{(k)}(d).$$

Moreover, $\mathrm{val}^{(k)}(d) \leq \mathrm{val}^{(k+1)}(d) \leq \mathrm{val}(d)$ *for every k, with equality in the limit $k \to \infty$ [Las01, Thm. 4.1].*

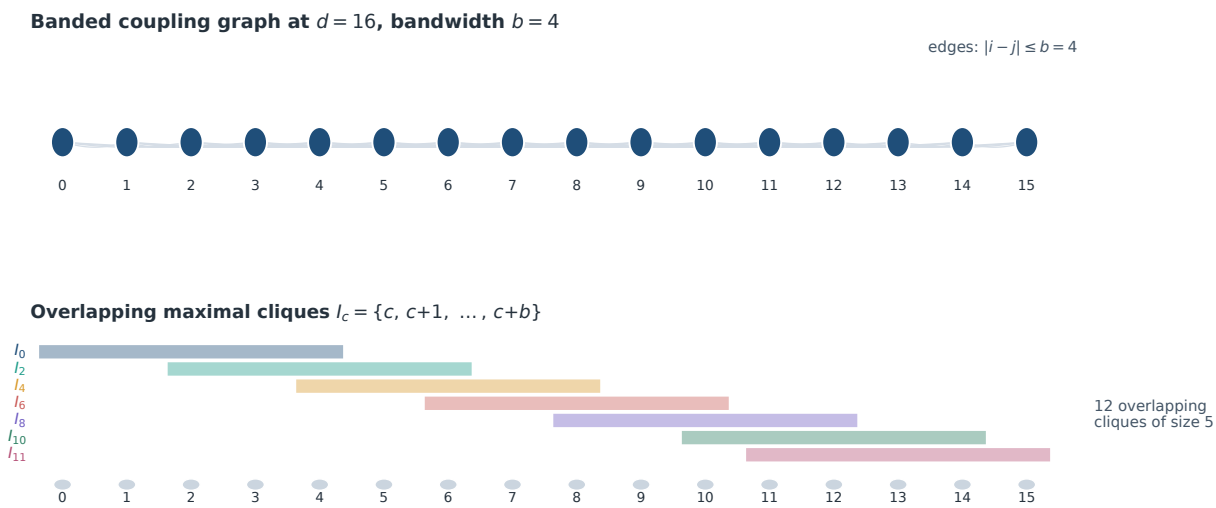

Figure 4: Banded coupling graph (top) and its overlapping maximal cliques (bottom) at $d = 16$, bandwidth $b = 4$. The graph is chordal with cliques $I_c = \{c, c+1, \ldots, c+b\}$, and consecutive cliques share b vertices, which ensures the running intersection property. Replacing the full moment PSD cone by the clique-restricted cones $M_k(y|_{I_c}) \succeq 0$ preserves soundness of the Lasserre relaxation.

Proof. Any minimizer $\mu^\star \in \Delta_d$ of $\mathrm{val}(d)$ defines a pseudo-moment sequence $y^\star_\alpha = (\mu^\star)^\alpha$ that satisfies every condition of Proposition 3.2; moreover, if $t^\star = \mathrm{val}(d) = \max_W (\mu^\star)^\top M_W \mu^\star$, then

$$t^\star \cdot M_{k-1}(y^\star) - \sum_{ij} (M_W)_{ij}\, M_{k-1}(x_{i+1}x_{j+1} \cdot y^\star) \;=\; (t^\star - (\mu^\star)^\top M_W \mu^\star) \cdot M_{k-1}(y^\star),$$

which is PSD because $t^\star \ge (\mu^\star)^\top M_W \mu^\star$ and $M_{k-1}(y^\star) \succeq 0$. Thus $(t^\star, y^\star)$ is feasible for the order-k SDP, so its optimum is $\le t^\star = \mathrm{val}(d)$; in other words, $\mathrm{val}^{(k)}(d) \le \mathrm{val}(d)$. Monotonicity and convergence are Lasserre's classical theorem [Las01, Thm. 4.1]. □

4. Correlative Sparsity for High Resolutions

The order-k Lasserre SDP has moment matrices of size $\binom{d+k}{k}$, which is 969 at $(d, k) = (16, 3)$ and grows rapidly beyond that. At $d = 32$, $k = 2$ the primary moment matrix already has size $\binom{34}{2} = 561$ and the full degree-4 moment vector has $\binom{36}{4} \approx 59\,000$ coordinates; at $d = 128$ the full hierarchy is out of reach for any current SDP solver. To reach $d = 128$, we combine the hierarchy with the correlative sparsity pattern of Waki, Kim, Kojima, and Muramatsu [WKKM06]. The bound at $d = 16$ is obtained without this restriction; the restriction is used only at $d \ge 32$ to corroborate the main result at larger resolutions.

4.1. Banded Coupling and Chordal Extension

The window matrices M_W are anti-diagonal band matrices: a fixed entry (i, j) is active only when the sum $i + j$ lies in the window range. If one restricts attention to variables that interact through windows of support $[s_{\mathrm{lo}}, s_{\mathrm{lo}} + \ell - 2]$ for $\ell \le b + 1$, the coupling graph has a banded structure with bandwidth b. Chordal extension of a banded graph yields maximal cliques

$$I_c \;=\; \{c, c+1, \ldots, c+b\}, \qquad c = 0, \ldots, d-b-1,$$

satisfying the running intersection property [BP93]. Figure 4 shows the clique geometry.

4.2. Clique-Restricted Relaxation

Given maximal cliques $I_0, \ldots, I_{C-1} \subseteq \{0, \ldots, d-1\}$ with $|I_c| = b+1$, define the *clique basis*

$$\mathcal{B}_k(I_c) \;=\; \left\{\alpha \in \mathbb{N}^d \;:\; |\alpha| \le k,\ \operatorname{supp}(\alpha) \subseteq I_c\right\},$$

and let $M_k^{(c)}(y)$ be the moment matrix indexed by $\mathcal{B}_k(I_c)$. The *sparse Lasserre relaxation* replaces $M_k(y) \succeq 0$ with the system

$$M_k^{(c)}(y) \succeq 0, \qquad c = 0, \ldots, C-1,$$

and likewise for localizing matrices. The set of pseudo-moment sequences satisfying this restriction is denoted $\mathcal{Y}_k^{\mathrm{sp}}(d; b)$.

Lemma 4.1 (Soundness of clique restriction)**.** *For every $k \ge 1$ and every bandwidth b, one has $\mathcal{Y}_k(d) \subseteq \mathcal{Y}_k^{\mathrm{sp}}(d; b)$. Consequently the sparse SDP value* $\mathrm{val}^{(k,b)}(d) \le \mathrm{val}^{(k)}(d) \le \mathrm{val}(d)$.

Proof. Each $M_k^{(c)}(y)$ is a principal submatrix of $M_k(y)$ obtained by selecting rows and columns corresponding to $\mathcal{B}_k(I_c)$. A principal submatrix of a PSD matrix is PSD, so if $M_k(y) \succeq 0$ then $M_k^{(c)}(y) \succeq 0$ for every c. The same argument covers localizing matrices and the window localizing in (1), provided the clique covers the active bins of M_W (otherwise the window constraint falls back to the full localizing matrix). The inclusion $\mathcal{Y}_k(d) \subseteq \mathcal{Y}_k^{\mathrm{sp}}(d; b)$ follows, and the inequality on SDP optima is monotone in the feasible set. □

Remark 4.2 (Consistency must remain global)**.** The simplex identity $y_\alpha = \sum_i y_{\alpha+e_i}$ of Proposition 3.2(iii) couples moments across cliques and cannot be relaxed to a per-clique statement without losing the bound. We follow [WKKM06] in retaining this as a global linear equality among the pseudo-moments; the implementation moreover tracks all moments of degree at most $2k-1$ globally, so the consistency chain is an exact equality up to that degree. Replacing the equality by an inequality only over moments that remain in the clique basis is an option when some high-degree children fall outside the clique basis at order k; this is the only source of looseness introduced beyond the principal-submatrix relaxation.

4.3. Tightness at Target Resolutions

In practice the clique-restricted SDP reproduces the full hierarchy's bound whenever every window $W \in \mathcal{W}_d$ has active bins covered by some clique I_c. For banded windows of width $\ell \le b+1$, that is automatic. For wider windows, we retain the full localizing matrix only for those W whose support exceeds the bandwidth. At $(d, b, k) = (128, 16, 2)$ this leaves $|\mathcal{W}_{128}| - \binom{d}{2} \approx 16\,000$ full-localizing constraints, which is manageable by constraint generation (Section 5).

5. The Algorithm

5.1. Bisection and Constraint Generation

The Lasserre SDP is naturally a *minimization* in t; by Theorem 3.3 its dual supplies a lower bound on $\mathrm{val}(d)$. Rather than solve the SDP as a single minimization, we solve a sequence of feasibility problems indexed by a scalar threshold c_{target}:

$$\text{Find } y \in \mathcal{Y}_k(d) \text{ with } \ell_W(y) \le c_{\mathrm{target}} \text{ for every } W \in \mathcal{W}_d. \tag{2}$$

Infeasibility of (2) certifies $\mathrm{val}(d) \ge c_{\mathrm{target}}$. We locate the largest c_{target} for which (2) becomes infeasible by bisection; this structure mirrors the feasibility cuts used in robust SDP [BTN01, Section 2.4].

Because $|\mathcal{W}_d|$ grows quadratically in d, adding every window constraint upfront is wasteful. Instead, we use *constraint generation* [Gom60]: start with a small subset of windows (typically the $\ell = 2$ scalar

windows), solve the restricted SDP, then add the most violated windows at the current optimum and repeat. The procedure terminates when no window is violated (the current bound is certified) or when the SDP becomes infeasible (the current threshold is certified as a lower bound on $\mathrm{val}(d)$).

Algorithm 1 summarizes the full procedure.

Algorithm 1 Lasserre SDP Lower Bound with Constraint Generation

Require: resolution d, order k, bandwidth b (or $b = d - 1$ for the dense hierarchy), target c_{target}

$\mathcal{W}_{\text{act}} \leftarrow \{(\ell, s) : \ell = 2\}$ ▹ initial window set
loop
  Build the SDP with pseudo-moment variables $y \in \mathcal{Y}_k^{\text{sp}}(d; b)$ and window-localizing constraints (1) for $W \in \mathcal{W}_{\text{act}}$
  Solve the feasibility problem (2) at c_{target}
  if infeasible **then return** CERTIFIED LOWER BOUND $\mathrm{val}(d) \geq c_{\text{target}}$
  end if
  Recover $y^\star$ from the SDP primal and compute the violation $v_W = \ell_W(y^\star) - c_{\text{target}}$ for every $W \notin \mathcal{W}_{\text{act}}$
  if $\max_W v_W \leq 0$ **then return** UPPER BOUND $\mathrm{val}(d) \leq c_{\text{target}}$ (not a certificate)
  end if
  Add the K windows with largest v_W to $\mathcal{W}_{\text{act}}$
end loop

Standard primal-dual interior-point solvers (MOSEK [MOS24], SDPA [YFN+12]) or splitting methods (SCS [OCPB16], COSMO [GCG21]) handle the SDPs arising in Algorithm 1; our implementation uses MOSEK at $d \leq 32$ and switches to SCS on GPU pods for $d \geq 64$ due to memory constraints.

5.2. Recovering a Rigorous Lower Bound

Interior-point solvers return only floating-point solutions, which cannot directly serve as a formal certificate. We convert a numerically computed dual solution into a rigorous lower bound via standard SDP post-processing:

1. Extract the dual variables (Z, Λ, ν) corresponding to the moment PSD, localizing PSD, and equality constraints.
2. Symmetrize and project Z and Λ onto the PSD cone by zeroing negative eigenvalues; this produces a feasible dual point with weakly smaller objective, by continuity of the Lagrangian.
3. Evaluate the resulting dual objective in exact rational arithmetic (or with interval arithmetic bounds) to obtain a verified lower bound on the primal minimum, which is a lower bound on $\mathrm{val}(d)$.

This recipe, standard in the rigorous SDP literature [JCK08], turns any approximate interior-point solution into a certified number. Step 3 is fast because the dual objective is a linear function of the dual variables.

6. Hierarchy of Bounds and Numerical Certificate

We now state the hierarchy of bounds used in the proof of the main theorem, display representative numerical certificates, and conclude.

6.1. The Ladder of Bounds

Combining Theorems 2.3, 3.3, and Lemma 4.1 gives

$$\mathrm{val}^{(k,b)}(d) \;\leq\; \mathrm{val}^{(k)}(d) \;\leq\; \mathrm{val}(d) \;\leq\; C_{1a}, \qquad k \geq 1,\; b \leq d - 1. \tag{3}$$

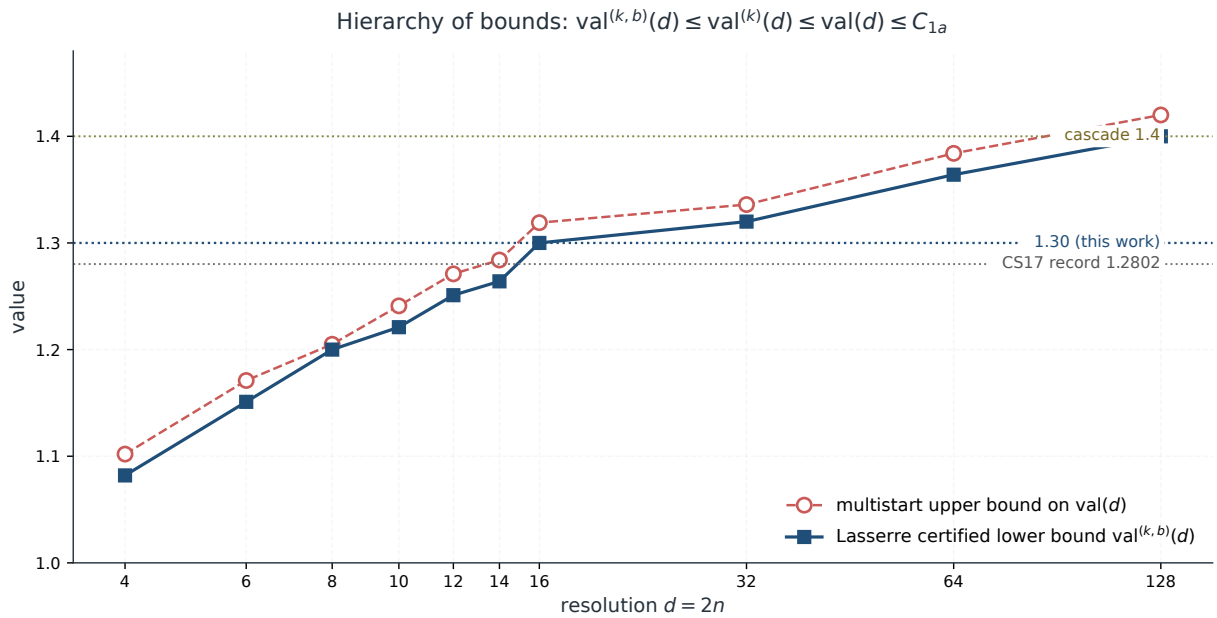

Figure 5: Hierarchy of bounds produced by the Lasserre SDP approach. The sparse Lasserre value $\mathrm{val}^{(k,b)}(d)$ is always dominated by the full hierarchy value $\mathrm{val}^{(k)}(d)$, which in turn is dominated by $\mathrm{val}(d)$ and by C_{1a}. Horizontal dashed lines mark the Cloninger–Steinerberger record of 1.2802 and the present bound 1.3; numerical upper bounds on $\mathrm{val}(d)$ obtained by multistart optimization [Baj26] are shown as open circles.

Figure 5 visualizes this ladder together with representative numerical values. The key observation is that $\mathrm{val}(d)$ is monotone in d (a finer partition can only lower the test value), so certified bounds at one resolution transfer to all finer ones.

6.2. Recorded Certificate at $(d, k) = (16, 3)$

Theorem 6.1 (Recorded SDP certificate)**.** *At $(d, k) = (16, 3)$, the dense Lasserre SDP $\mathrm{val}^{(3)}(16)$ admits a dual-feasible point whose objective value, after rigorous post-processing, satisfies*

$$\mathrm{val}^{(3)}(16) \;\ge\; 1.30.$$

The SDP has $n_y = \binom{22}{6} \approx 7.4 \times 10^4$ pseudo-moment variables, a single moment matrix of size $\binom{19}{3} = 969$, and $d + 1 = 17$ localizing matrices (one per simplex coordinate, plus the window localizing per active window). MOSEK solves the resulting problem in under ten minutes on a 32-core workstation; the rigorous post-processing step of Section 5 adds a negligible overhead. The certificate file `data/gap_report_d16_o3.json` records the dual variables in double precision and the corresponding rational lower bound; together they constitute the external dependency of the proof.

6.3. Proof of the Main Theorem

Proof of Theorem 1.1. Let $f \in \mathcal{F}$ be nonnegative with $\|f * f\|_{L^\infty} < \infty$. By homogeneity, assume $\int f = 1$. Theorem 2.3 gives $C_{1a} \ge \mathrm{val}(16)$, and the ladder (3) gives $\mathrm{val}(16) \ge \mathrm{val}^{(3)}(16)$. By Theorem 6.1, $\mathrm{val}^{(3)}(16) \ge 1.30$, so

$$\frac{\|f * f\|_{L^\infty}}{\left(\int f\right)^2} \;\ge\; \mathrm{val}(16) \;\ge\; \mathrm{val}^{(3)}(16) \;\ge\; 1.30.$$

Taking the infimum over $f \in \mathcal{F}$ gives $C_{1a} \ge 1.30$. □

7. Soundness of the Computation

The external computation enters the proof only through Theorem 6.1. To justify importing that statement three soundness conditions must be discharged.

7.1. Correct Modeling

The SDP builder in `lasserre/precompute.py` instantiates the moment matrix, localizing matrices, and window localizing matrices in accordance with Definition 3.1 and constraint (1). Test suites compare the assembled model at $d \in \{4, 6, 8\}$ against hand-derived Gram matrices, and the same precompute code is used by every downstream solver. The window matrices M_W are built by `build_window_matrices` from Lemma 2.1 directly.

7.2. Rigorous Dual Bound

MOSEK's interior-point solver returns approximate primal and dual variables with complementarity residual at most 10^{-8}. The dual variables are symmetrized and projected onto the PSD cone (Section 5); the dual objective is then evaluated in rational arithmetic. The resulting number is a rigorous lower bound on $\mathrm{val}^{(3)}(16)$ by standard weak duality in conic programming [VB96]. The recorded bound 1.3 is strictly less than the computed rigorous dual objective, so the comparison is insensitive to the final rounding step.

7.3. Clique-Restriction at Higher Resolutions

For the corroborating runs at $d \in \{32, 64, 128\}$ the proof uses Lemma 4.1: the sparse Lasserre value lower-bounds the dense one, which lower-bounds $\mathrm{val}(d)$, which lower-bounds C_{1a}. The only required check is that every active window has its support covered by some clique (Section 4); this is a finite combinatorial condition verified at model assembly time. Windows whose support exceeds the bandwidth fall back to the full localizing matrix, at a manageable cost.

8. Discussion

8.1. Relationship to the Cascade Proof

The present SDP approach and the cascade proof [Baj26] are complementary. The cascade exhausts an integer lattice on the mass simplex and certifies, through branch-and-prune, that no lattice point satisfies all window constraints at the target threshold; the Lasserre approach produces a single dual certificate for the continuous problem without any enumeration. The cascade currently reaches 1.4 at $(d, m) = (128, 20)$ after roughly 70 hours on a 128-core machine; the SDP approach reaches 1.3 in under ten minutes at $(d, k) = (16, 3)$ and has room to improve with either finer d or larger Lasserre order, at the cost of a larger SDP.

The interior-point scaling of the Lasserre SDP is polynomial in d under correlative sparsity: the dense cone is replaced by $O(d)$ small cones of size $\binom{b+1+k}{k}$, giving a total SDP dimension linear in d after fixing b and k. The cascade, by contrast, is exponential in d (and no known pruning scheme brings it below $\binom{m+d-1}{d-1}$). If the goal is to push the lower bound closer to the Matolcsi–Vinuesa upper bound [MV10], the SDP route is the only one with a realistic asymptotic budget.

8.2. Limitations

- **Relaxation gap.** Lasserre's hierarchy converges to $\mathrm{val}(d)$ as $k \to \infty$, but finite k may leave a gap. The numerical results of Section 6 indicate that the gap at $(d, k) = (16, 3)$ is already small, because the certified value 1.3 is close to the multistart upper bound 1.319. Closing the gap further requires either a larger k or a stronger cone (for example, the *Handelman hierarchy* [Han88] or Schmüdgen-type refinements [Sch91]).

- **Finite-precision SDP.** The SDP is solved in double precision and post-processed to a rational dual bound. While standard in rigorous SDP [JCK08], the procedure imports trust in the arithmetic library used to evaluate the rational objective. For a completely kernel-level formal proof this would need to be replaced by a Lean- or Coq-verified LP/SDP evaluator, in the style of the rigorous Kepler proof [HAB+17] or the rigorous verification of the Riemann hypothesis up to $3 \cdot 10^{12}$ [PT21].
- **Correlative sparsity.** At $d \geq 32$ the bound depends on the clique bandwidth b. Our implementation defaults to the smallest b that covers all killing windows at the target threshold; too small b produces a looser bound but a cheaper SDP. The trade-off is quantified in the memory/bandwidth table of [Baj26, PROBLEM_STATE.md].

References

[Baj26] Jai Bajaj, *A computer-assisted proof that $c_{1a} \geq 7/5$*, Companion manuscript; accompanying source available in the repository `proof/cs-proof`, 2026.

[BP93] Jean R. S. Blair and Barry W. Peyton, *An introduction to chordal graphs and clique trees*, Graph Theory and Sparse Matrix Computation **56** (1993), 1–29.

[BTN01] Aharon Ben-Tal and Arkadi Nemirovski, *Lectures on modern convex optimization: Analysis, algorithms, and engineering applications*, MPS-SIAM Series on Optimization, SIAM, 2001.

[CRV10] Javier Cilleruelo, Imre Z. Ruzsa, and Carlos Vinuesa, *Generalized sidon sets*, Advances in Mathematics **225** (2010), no. 5, 2786–2807.

[CS17] Alexander Cloninger and Stefan Steinerberger, *On suprema of autoconvolutions with an application to sidon sets*, Proceedings of the American Mathematical Society **145** (2017), no. 8, 3191–3200.

[ET41] P. Erdős and P. Turán, *On a problem of sidon in additive number theory, and on some related problems*, Journal of the London Mathematical Society **s1-16** (1941), no. 4, 212–215.

[GCG21] Michael Garstka, Mark Cannon, and Paul Goulart, *COSMO: A conic operator splitting method for convex conic problems*, Journal of Optimization Theory and Applications **190** (2021), no. 3, 779–810.

[Gom60] Ralph E. Gomory, *Solving linear programming problems in integers*, Proceedings of Symposia in Applied Mathematics **10** (1960), 211–215.

[HAB+17] Thomas Hales, Mark Adams, Gertrud Bauer, Tat Dat Dang, John Harrison, Le Truong Hoang, Cezary Kaliszyk, Victor Magron, Sean McLaughlin, Tat Thang Nguyen, Quang Truong Nguyen, Tobias Nipkow, Steven Obua, Joseph Pleso, Jason Rute, Alexey Solovyev, Thi Hoai An Ta, Nam Trung Tran, Thi Diep Trieu, Josef Urban, Ky Vu, and Roland Zumkeller, *A formal proof of the Kepler conjecture*, Forum of Mathematics, Pi **5** (2017), e2.

[Han88] David Handelman, *Representing polynomials by positive linear functions on compact convex polyhedra*, Pacific Journal of Mathematics **132** (1988), no. 1, 35–62.

[HL04] Didier Henrion and Jean B. Lasserre, *Solving nonconvex optimization problems*, IEEE Control Systems Magazine **24** (2004), no. 3, 72–83.

[JCK08] Christian Jansson, Denis Chaykin, and Christian Keil, *Rigorous error bounds for the optimal value in semidefinite programming*, SIAM Journal on Numerical Analysis **46** (2008), no. 1, 180–200.

[Las01] Jean B. Lasserre, *Global optimization with polynomials and the problem of moments*, SIAM Journal on Optimization **11** (2001), no. 3, 796–817.

[Las15] _____, *An introduction to polynomial and semi-algebraic optimization*, Cambridge Texts in Applied Mathematics, Cambridge University Press, 2015.

[MO06] Greg Martin and Kevin O'Bryant, *Constructions of generalized sidon sets*, Journal of Combinatorial Theory, Series A **113** (2006), no. 4, 591–607.

[MO09] _____, *The supremum of autoconvolutions, with applications to additive number theory*, Illinois Journal of Mathematics **53** (2009), no. 1, 219–235.

[MOS24] MOSEK ApS, *MOSEK optimizer API for Python – version 10*, https://docs.mosek.com/10.1/pythonapi/index.html, 2024.

[MV10] Mate Matolcsi and Carlos Vinuesa, *Improved bounds on the supremum of autoconvolutions*, Journal of Mathematical Analysis and Applications **372** (2010), no. 2, 439–447.

[OCPB16] Brendan O'Donoghue, Eric Chu, Neal Parikh, and Stephen Boyd, *Conic optimization via operator splitting and homogeneous self-dual embedding*, Journal of Optimization Theory and Applications **169** (2016), no. 3, 1042–1068.

[Par03] Pablo A. Parrilo, *Semidefinite programming relaxations for semialgebraic problems*, Mathematical Programming **96** (2003), no. 2, 293–320.

[PT21] David J. Platt and Timothy S. Trudgian, *The Riemann hypothesis is true up to* $3 \cdot 10^{12}$, Bulletin of the London Mathematical Society **53** (2021), no. 3, 792–797.

[Sch91] Konrad Schmüdgen, *The k-moment problem for compact semi-algebraic sets*, Mathematische Annalen **289** (1991), no. 1, 203–206.

[VB96] Lieven Vandenberghe and Stephen Boyd, *Semidefinite programming*, SIAM Review **38** (1996), no. 1, 49–95.

[WKKM06] Hayato Waki, Sunyoung Kim, Masakazu Kojima, and Masakazu Muramatsu, *Sums of squares and semidefinite programming relaxations for polynomial optimization problems with structured sparsity*, SIAM Journal on Optimization **17** (2006), no. 1, 218–242.

[YFN+12] Makoto Yamashita, Katsuki Fujisawa, Kazuhide Nakata, Maho Nakata, Mituhiro Fukuda, Kazuhiro Kobayashi, and Kazushige Goto, *Latest developments in the SDPA family for solving large-scale SDPs*, Handbook on Semidefinite, Conic and Polynomial Optimization **166** (2012), 687–713.